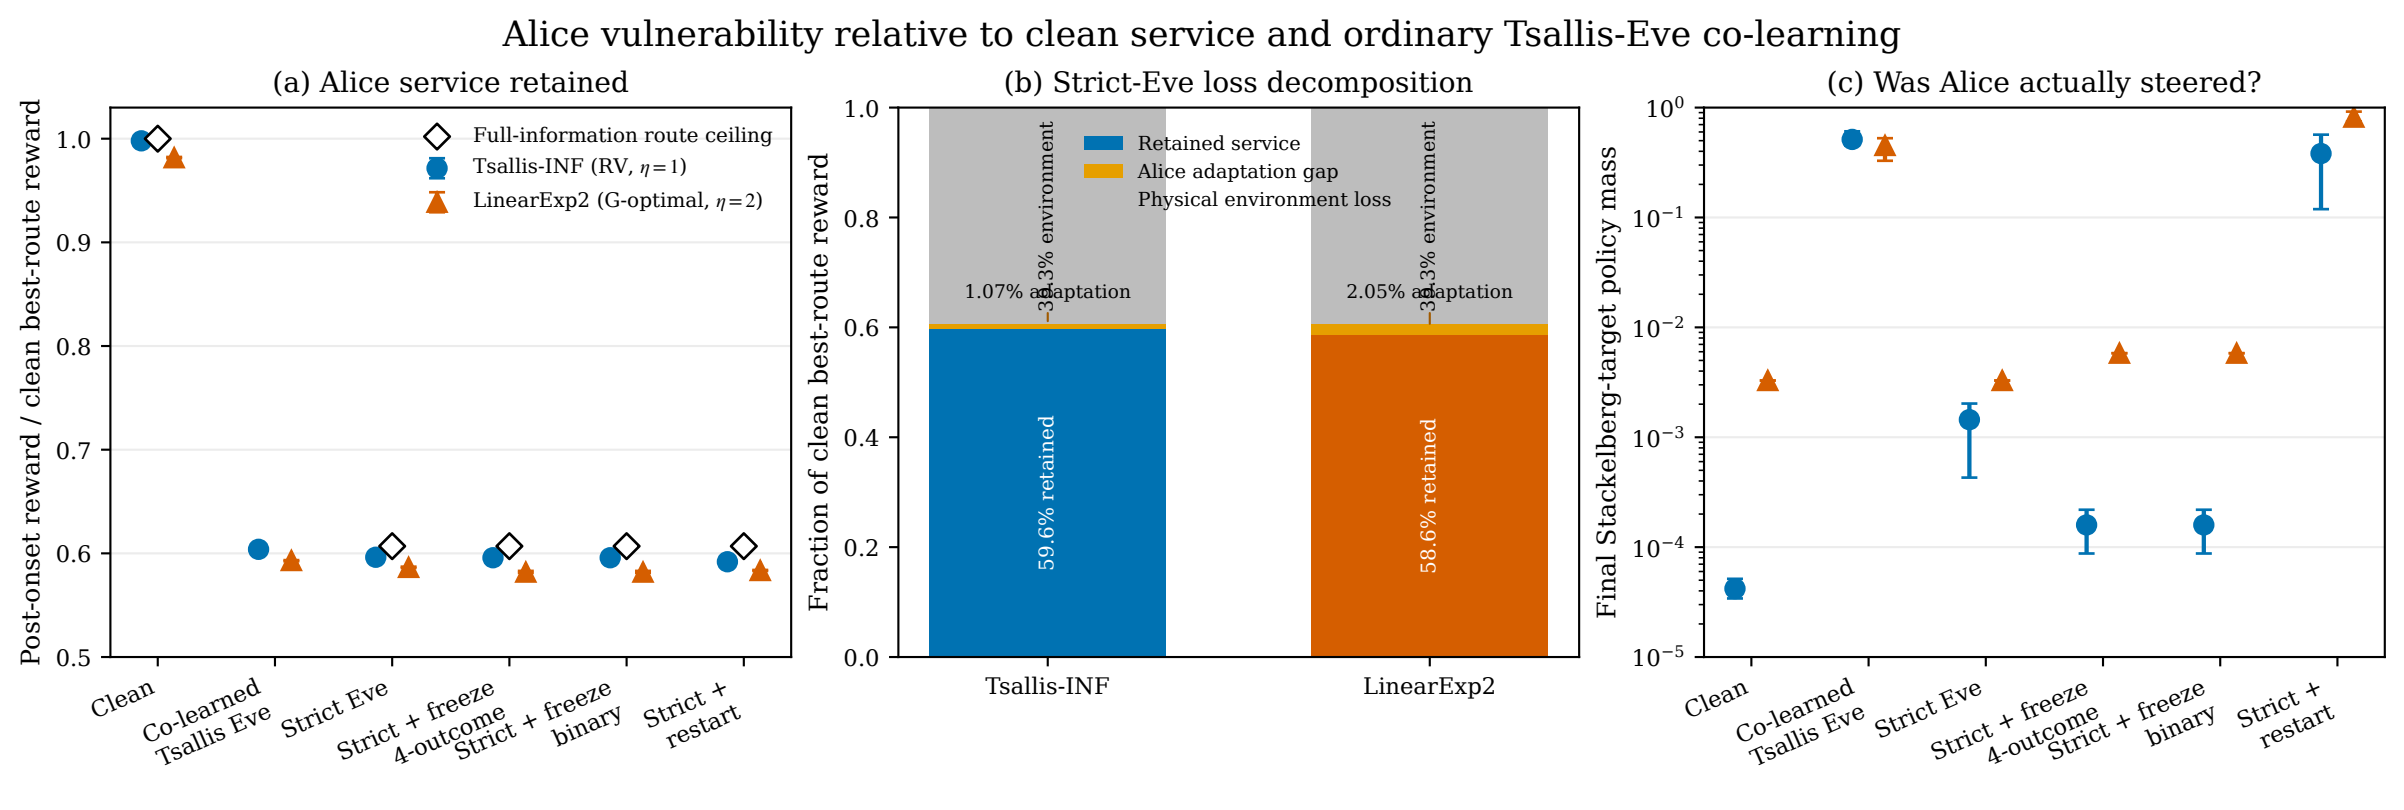
Alice vulnerability relative to clean service and ordinary Tsallis-Eve co-learning
(a) Alice service retained
Post-onset reward / clean best-route reward
Full-information route ceiling
Tsallis-INF (RV, η=1)
LinearExp2 (G-optimal, η=2)
Clean
Co-learned Tsallis Eve
Strict Eve
Strict + freeze 4-outcome
Strict + freeze binary
Strict + restart
(b) Strict-Eve loss decomposition
Fraction of clean best-route reward
Retained service
Alice adaptation gap
Physical environment loss
1.07% adaptation
33.3% environment
59.6% retained
2.05% adaptation
33.3% environment
58.6% retained
Tsallis-INF
LinearExp2
(c) Was Alice actually steered?
Final Stackelberg-target policy mass
Clean
Co-learned Tsallis Eve
Strict Eve
Strict + freeze 4-outcome
Strict + freeze binary
Strict + restart